

Annotation Guidelines

General Hints

- Please always annotate a tweet in isolation. You should never take into consideration tweets preceding or following the tweet to be annotated. For example, given the two consecutive tweets from a particular timeline:

- My girlfriend @USER_i is coming over this afternoon
- Apart from Jack and me, @USER_i and her new boyfriend will be there

This second tweet should be annotated as being biographically irrelevant. This is so, since one can only infer its relevance by considering the knowledge of the first tweet (which establishes that @USER_i is the author's girlfriend).

If tweets contain links or URLs, then the respective external sources should not be considered for annotation. Your annotation should exclusively focus on the text which is immediately contained in the tweet itself. This does not include contents that are linked to the tweet. For instance, in the following example, from the tweet alone, one does not know which song the author likes because this information resides on the website linked via the given URL:

- Love love love this song. @USER URL

- Twitter is a social-media platform in which predominantly informal language is employed. Vulgar language, such *shit* or *bitch*, occurs quite often. However, most of the time, it is used as an intensifier rather than as a means to offend anyone or anything.

- I'd do the same shit wtf URL

The above tweet, for instance, does not convey that its author dislikes the thing intensely that they are talking about. After all, they would consider doing it themselves.

If a tweet contains many common swearwords, then this should be regarded as a sign of *street language*. In the following tweet, for instance, the b-word is not meant in an insulting way. Neither should the n-word (in that tweet) be interpreted as some form of racism:

- I wanna know how ya bitches that never had ya own nigga be feeling fr?

As a consequence, such cases should not be regarded as some form of biographical relevance. However, this is different in examples in which such swearwords are meant to insult somebody:

- I know. FUCK DONALD TRUMP @USER

Category Inventory

The following category inventory is more fine-grained than the binary classification scheme that we presented in the paper and subsequently used in our supervised learning experiments. We describe this fine-grained scheme here since it is the scheme that was used for the actual annotation. Our final gold standard with its coarse-grained categories was produced in the following way:

- Instances labeled as **Biographically Relevant** (fine-grained category!) and **Unambiguously Biographically Relevant but Inferred** were merged to the coarse-grained category *BIOGRAPHICALLY RELEVANT*.

- Instances labeled as **Biographically Irrelevant** (fine-grained category!), **Ambiguously Biographically Relevant but Inferred** and **Missing Context** were merged to the coarse-grained category *BIOGRAPHICALLY IRRELEVANT*.
- Instances labelled as **Unclear/Uncertain** and **No 1st Person Tweet** were not included in the final dataset.

Fine-Grained Category: Biographically Relevant

All the information one would commonly write into a biography or a (user) profile are considered biographically relevant. This includes characteristics of a person or regular activities that a person participates in. It also includes one-time events if they represent major life time events (e.g. wedding, graduation or relocation).

Some common biographically relevant relations are: age, date of birth, place of birth, place of residence, age, job, physical characteristics (size, weight, eye colour, etc.), family status, family members and their relation to each other, hobbies, likes, dislikes, habits, personal traits, nationality, pets, major life events.

- Here are some prototypical examples:
 - Dumbest shit I've heard, my dad has always hung an American flag outside our house since I was little, cause we are PROUD TO BE AMERICANS , we ain't racist GTFO http
 - @USER I just moved from California to Augusta and all I can think about is @USER .. I wanna go to the sub station
 - I always smoke before I eat on lunch it makes eating so much better lol
 - I CANT STAND MY DAD AND HIS GIRLFRIEND SOME TIMES HOLY SHIIIT
 - One thing that irks the shit outta me is if a place doesn't have harmonious furniture/decor/order/colors
- Not only do we consider regular activities that characterize a person (in some way) as biographically relevant but also regular occurrences that have a greater impact on the life of the pertaining person:
 - Why can't I ever sleep
- Having an account on some social-media platform other than Twitter¹ even if the respective user name on that platform is not specified is considered biographically relevant:
 - blocked this boy from my Facebook cause he reacted to everybody happy birthday comment but mine
- We also include characteristics that are based on astrology as long as they are not episodic:
 - @USER my uranus is in Aquarius yes
- Assessments and evaluations are included. They may be highly subjective and do not have to correspond to the (objective) truth:
 - @USER I'm ugly save your camera

¹ Remember, you are annotating tweets for Twitter, so each author is already known to have an account on Twitter. Therefore, only the remaining social-media platforms are relevant.

- Information on possessions are included if these are possessions that not (nearly) every person has and whose property has a larger impact on the life of the person who owns it (e.g. car, garden, yacht):
 - My backyard is so green
 - I couldn't get my car fixed I'm looking for a new one

Fine-Grained Category: Biographically Irrelevant

We consider the following cases as biographically irrelevant tweets:

- Tweets that tell us (virtually) nothing about the author of a tweet but focus on a different topic:
 - I guess they're also too cool to wear a coat. URL
 - @USER I bet those two made a wicked pasta dinner. URL
- Characteristics or regular activities that apply to (virtually) all people:
 - I hate feeling down for no reason
 - I hate missing people I'm not supposed to miss
- One-off activities or events that do not represent major life-time events or imply anything special about the author of the tweet:
 - And i just got tabasco in my eye
 - @USER I entered an incorrect address for a order and I am trying to change it before it ship. I can't figure out how to change the shipping address or who I can contact.
- Personal trivia that ultimately tell us nothing special about the author of the tweet (i.e. contents that would never make it into a biography or user profile):
 - How long it take y'all to drink a caprisun? It takes me 5 seconds.
 - @USER Never once have I met a white Gilbert
 - Why do I always hear shit in my house when I'm home alone ???

Many possessions should be counted as trivia:

- in case you were wondering i am now the proud owner of a STRAWBERRYICEY BUCKET HAT my new all time favorite purchase omfg I'mnever taking it off @USER on IG & Etsy!!!! URL

One should decide on the basis of the impact on the life reality: if the possession of some item has a large impact, then this should be regarded as biographically relevant. Please also pay attention that the possession may be an unambiguous hint of some biographically information. In that case, you should categorize the utterance as biographically relevance. If you need to infer that biographical information and this information may not be guaranteed (for example, the possession of a swimsuit may imply the ability to swim but this is not for certain), then you should use other categories that are described in the remainder of these guidelines. Only if a possession has no significant impact on the author's life and there is not even an ambiguous hint of some biographical information, then you should use the present category (i.e. **Biographically Irrelevant**).

- Sympathy/Antipathy towards other users on Twitter or persons count as biographically irrelevant as long as these persons are not well-known (otherwise we have an instance of biographical relevance) and there is nothing known about the person that may imply anything biographically relevant about the author of the tweet:
 - @USER kyla!!! I'm so happy to see you back on the tl YOU'RE amazing

Fine-Grained Category: Unambiguously Biographically Relevant but Inferred

The tweet conveys some biographically relevant information but in an indirect way:

- @USER i too am drinking wine cheers you lovely human

Explanation: Drinking wine is not biographically relevant information but we may infer from this that the author is drinking alcohol, in general, which is regarded as biographically relevant information. This inference is also unambiguous.

- uhm any book recommendations? i see a day trip to barnes n noble in my near future for mental health care

Explanation: This tweet implies that the author is a keen reader.

- I woke up this morning, feeling terrible! It's okay because I know that God has me in his arms!

Explanation: This tweet implies that the author is a religious person.

- My grandmother tryna give me a speech about life just cause I said my sneakers cost 650\$ 🤔🤔

Explanation: From that tweet we can infer that the author's grandmother is still alive. This narrows down the author's age range (i.e. they must be a child or young adult).

- This category is, of course, related to the category **Biographical Relevant** and the following category **Ambiguously Biographically Relevant but Inferred**. The boundaries may not always be clear-cut.

Fine-Grained Category: Ambiguously Biographically Relevant but Inferred

One may infer some biographically relevant information from a tweet but one may not be sure that this inference is actually true. In other words, the tweet is ambiguous.

- Omw to church. :)

Explanation: If someone is on their way to church, this may be regarded as a clue that they are a religious person. However, there may be other possible reasons for visiting a church (e.g. being invited to a wedding).

- o Everybody round me all like I love my Bae we doing this & that for Christmas & I'm like well yeah
looks down those are some nice shoes

Explanation: The tweet sounds as if the author is a single or, at least, has relationship problems. However, this is not explicitly stated and one may only come to such conclusion since we try to find some explanation for the author's behaviour.

Fine-Grained Category: Missing Context

This category comprises tweets that may imply some biographical relevance but some context is missing in order to understand the entire statement. Examples are:

- The entity or activity which is referred to by some pronoun or adverb is not mentioned

- o We love that too URL

Explanation: Information that resides on some webpage linked to the tweet by some URL is considered missing context.

- o Thisssss is the way I live

Explanation: For this category, there needs to be some significant evidence for biographical relevance in the tweet itself (e.g. likes/dislikes or habits) for biographical relevance. Therefore, tweets that contain pronouns whose antecedent is missing and whose resolution would likely reveal no biographically relevant information would not be an instance of **Missing Context**, such as the following tweet:

- @USER He's ignored me, active 10 minutes ago
:/

(This tweet is biographically irrelevant.)

- The tweet looks as if it contained biographically relevant information but there is some ambiguity which needs to be resolved in order to specify the biographically relevant information. (In other words, the information that is missing is not that important for deciding whether the utterance is biographically relevant or biographically irrelevant.)

- o I think it was cause I didn't go to class!

Explanation: The author of the tweet seems to visit some form of class. However, it is not clear whether this is a class taught at school, at some university or some other educational institution (all of these alternatives would be considered biographically relevant information). Further, it is not clear when in the past this class was taught.

- o Its been 1 year exactly since ive had my licence

Explanation: It is unclear what type of licence the author possesses. No matter what type of licence it is, it would count as biographical relevance.

- Tweets that look as if they represent an answer to a specific but missing question revealing some biographically relevant information do not count as **Missing Context**:

- o @USER me.

Explanation: Although this tweet could represent some biographically relevant information, if it were the answer to a tweet, such as Does anyone live in Kansas and play the guitar really well? We're looking for

someone to join our band, this answer is to be labelled as **Biographically Irrelevant** since this tweet could similarly be a reply to an unimportant question.

o Ooou we got 6

Explanation: If this tweet were the answer to a question, such as How many children have you got?, the tweet would be biographically relevant. However, this tweet is to be labelled as **Biographically Irrelevant** since there could also be many irrelevant contexts possible (for instance, if the author had bought 6 onions, this would not count as biographically relevant information).

Fine-Grained Category: Unclear/Uncertain

This category should be labelled if:

- One cannot understand the tweet.
- One is completely uncertain about the actual category of the tweet.

Fine-Grained Category: No 1st Person Tweet

This category should be used if:

- The tweet has mistakenly been identified as a 1st person tweet (e.g. imperatives):
 - o Don't doooooo thatExplanation: We also extract tweets that do not contain a 1st person pronoun but begin with a verb since these are usually instances of *implicit protagonists* (e.g. joined a gym today)
- The tweet (or parts of it) is written in a language other than English:
 - o @USER @USER @USER Je me posais la question justement, on en produit tellement que l'on ne sais plus ou les mettre.
 - o I think it's funny how some broads really think we out here sweating them... oye mami no te hages falsas ilusiones que no se te ve bien .🤔🙄
- The tweet contains a 1st person pronoun but it does not refer to the author of the tweet. This can happen in direct speech:
 - o Guys in Augusta, GA be like "I know a spot" and take you here URL

or in lyrics:

- o Four in the morning I cannot sleep I am depressed sitting up in my bed The demons are roaring they're bothering me I'm sitting here with a gun

to my head Finger on trigger, what if I squeezed?
Does that mean that my problems will leave? Get
rid of what's bothering me?